

ASSIGNMENT 1
10 Different Model Providers and their Models released last 12 months

1. Provider : DeepSeek AI Agent: DeepSeek

Model	Release Date	Cut Off Date (Knowledge)	Token Costs (per Million tokens) and Context Window (CW)	Model Type	Usage Reasoning
R1-0528	May 28, 2025	Mar 31, 2025	Input: 0.50\$/M Output: 2.15\$/M CW: 163K Tokens	Fully Open Source public model	Provides enhanced reasoning performance and more accuracy (compared to initial R1 model). Significant reduction to hallucination rates
V3.1 and V3.1 Terminus	Aug 21, 2025 and Sep 22 2025	Mar 31, 2025	Input: 0.15\$/M Output: 0.75\$/M CW: 32K - 128K	Fully Open Source public model	A hybrid model which provides easy chat prompt switching between Thinking (Complex Reasoning) mode and Non-Thinking (Casual Q/A)
V3.2	Dec 1, 2025	May, 2025	Input: 0.252\$/M Output: 0.378\$/M CW: 131K	Fully Open Source public model	A hybrid model with Specialized reasoning
V4 Pro	Apr 24, 2026	May, 2025	Input: 0.435\$/M Output: 0.87\$/M CW: 1M	Private proprietary model	Provides higher efficiency and more Cost effective APIs with high throughput and strong reasoning
V4 Flash	Apr 24, 2026	May, 2025	Input: 0.14\$/M Output: 0.28\$/M CW: 1M	Fully Open Source public model	Provides higher efficiency and more Cost effective APIs with high throughput and strong reasoning

ASSIGNMENT 1
10 Different Model Providers and their Models released last 12 months

2. Provider: OpenAI Agent: GPT

Model	Release Date	Cut Off Date (Knowledge)	Token Costs (per Million tokens) and Context Window (CW)	Model Type	Usage Reasoning
GPT OSS	Aug 5, 2025	Jun 30, 2024	I/P: 0.039\$/M O/P: 0.18\$/M CW: 131K	Open Source/ Weight	Provides high-performance reasoning, agentic workflows, and local, private, or cost-effective inference
GPT 5.4	Mar 5, 2026	Aug 31, 2025	I/P: 2.50\$/M O/P: 15\$/M CW: 1M	Private Model	Provides Enhanced Reasoning & Coding by delivering superior performance in complex coding, debugging across multiple files, and understanding large codebases.
GPT 5.5	Apr 24, 2026	Dec 1, 2025	I/P: 5\$/M O/P: 30\$/M CW: 1M	Private Model	Specialized Agentic Model which provides stronger reasoning, higher reliability and improved token efficiency
GPT 5.3 Codex	Feb 24, 2026	Aug 31, 2025	I/P: 1.75\$/M O/P: 14\$/M CW: 400K	Private Model	Optimized for agentic coding and autonomous development of games and apps

ASSIGNMENT 1
10 Different Model Providers and their Models released last 12 months

3. Provider: Llama by Meta AI

Model	Release Date	Cut Off Date (Knowledge)	Token Costs (per Million tokens) and Context Window (CW)	Model Type	Usage Reasoning
Llama 4 Scout	Apr 6, 2025	Aug 31, 2024	I/P: 0.08\$/M O/P: 0.30\$/M CW: 1M	Open Source public Model	It is an MultiModal MoE excelling in long-document analysis, legal/finance summaries, and agentic tasks, RAG.
Llama 4 Maverick	Apr 6, 2025	Aug 31, 2024	I/P: 0.15\$/M O/P: 0.60\$/M CW: 1M	Open Source public Model	It is an MultiModal MoE excelling in superior reasoning, high-accuracy analysis and complex single-domain tasks.
Llama 4 Guard 12B	Apr 30, 2025	Aug 31, 2024	I/P: 0.18\$/M O/P: 0.18\$/M CW: 163K	Open Source public Model	offers superior safety moderation by natively supporting multimodal inputs (text and images), providing enhanced, high-performance content filtering designed specifically for Llama 4 models.

ASSIGNMENT 1
10 Different Model Providers and their Models released last 12 months

4. Provider: Google Agent : Gemini

Model	Release Date	Cut Off Date (Knowledge)	Token Costs (per Million tokens) and Context Window (CW)	Model Type	Usage Reasoning
Gemini 2.5 Flash lite	Jul 22, 2025	Jan 31, 2025	I/P: 0.10\$/M O/P: 0.40\$/M CW: 1M	Private Model	A highly optimized, lower-cost model designed for tasks needing low latency, such as classification and translation.
Gemini 2.5 Flash	Jun 17, 2025	Jan 31, 2025	I/P: 0.30\$/M O/P: 2.50\$/M CW: 1M	Private Model	The model is designed for advanced reasoning, coding, mathematics, and scientific tasks.
Gemini 2.5 Pro	Jun 17, 2025	Jan 31, 2025	I/P: 1.25\$/M O/P: 10\$/M CW: 1M	Private Model	The model is designed for advanced reasoning, coding, mathematics, and scientific tasks.
Gemini 3 Flash Preview	Dec 17, 2025	Jan 31, 2025	I/P: 0.50\$/M O/P: 3\$/M CW: 1M	Private Model	Provides configurable reasoning via thinking levels (minimal, low, medium, high), structured output, tool use, and automatic context caching, optimized for users who want strong reasoning and agentic behavior without the cost or latency of full scale frontier models.

ASSIGNMENT 1
10 Different Model Providers and their Models released last 12 months

Gemini 3.1 Pro Preview	Feb 19, 2025	Jan 31, 2025	I/P: 2\$/M O/P: 12\$/M CW: 1M	Private Model	It excels in complex reasoning, deep analysis, and multimodal tasks requiring high accuracy thus it is best for intricate coding or research.
-------------------------------	---------------------	---------------------	--	----------------------	--

5. Provider: Anthropic Agent: Claude

Model	Release Date	Cut Off Date (Knowledge)	Token Costs (per Million tokens) and Context Window (CW)	Model Type	Usage Reasoning
Claude Sonnet 4	Jan 31, 2025	Jan 31, 2025	I/P: 3\$/M O/P: 15\$/M CW: 1M	Private Model	Optimized for practical everyday use, providing advanced reasoning capabilities while maintaining efficiency and responsiveness in diverse internal and external scenarios.
Claude Opus 4	Aug 5, 2025	Jan 31, 2025	I/P: 15\$/M O/P: 75\$/M CW: 200K	Private Model	Supports extended, agentic workflows, handling thousands of task steps continuously for hours without degradation..
Claude Opus 4.1	Aug 5, 2025	Jan 31, 2025	I/P: 15\$/M O/P: 75\$/M CW: 200K	Private Model	Optimized for tasks involving research, data analysis, and tool-assisted reasoning.
Claude Opus 4.7	Apr 16, 2026	Jan, 2026	I/P: 5\$/M O/P: 25\$/M CW: 1M	Private Model	Built for long-running, asynchronous agents. effective for asynchronous agent pipelines where tasks unfold over time - large

ASSIGNMENT 1
10 Different Model Providers and their Models released last 12 months

					codebases, multi-stage debugging, and end-to-end project orchestration
Claude Haiku 4.5	Oct 15, 2025	Jan 31, 2025	I/P: 1\$/M O/P: 5\$/M CW: 200K	Private Model	Delivers near-frontier intelligence to real-time and high-volume applications. at a fraction of the cost and latency of larger Claude models.
Claude Sonnet 4.6	Feb 17, 2026	Jan, 2026	I/P: 3\$/M O/P: 15\$/M CW: 1M	Private Model	It excels at iterative development, complex codebase navigation, end-to-end project management with memory, polished document creation, and confident computer use for web QA and workflow automation.

6. Provider: Alibaba Cloud Agent: Qwen

Model	Release Date	Cut Off Date (Knowledge)	Token Costs (per Million tokens) and Context Window (CW)	Model Type	Usage Reasoning
Qwen 3.6 Series (Plus, Max Preview, Flash, 27B, 35B A3B)	Apr, 2026	Early 2026	I/P: 0.15\$ - 0.325\$/M O/P: 1 - 3.20\$/M CW: 262K - 1M	mixed Model(Plus is private, others are Open Source)	Series focuses on "vibe coding" and real-world agent tasks
Qwen 3.5	Feb & Apr,	Early 2026	I/P: 0.065\$ -	Open	A high-performance,

ASSIGNMENT 1
10 Different Model Providers and their Models released last 12 months

Series (Plus, 27B, 397B-A17B)	2026		0.40\$/M O/P: 0.15\$ - 2.40\$/M CW: 262K - 1M	Source	open-weight AI model designed as a native multimodal agent with significant advantages in efficiency and speed. Its core strengths include a sparse mixture-of-experts (MoE) architecture ,native vision/voice capabilities, and exceptional cost-efficiency
Qwen 3 Series (235B-A22B, 30B-A3B)	Apr 28, 2026	Mar 31 2025	I/P: 0.455\$/M O/P: 1.82\$/M CW: 131K	Open Source Public	Seamless switching between a "thinking" mode for complex reasoning, math, and code tasks, and a "non-thinking" mode for general conversational efficiency and also multilingual support

7. Provider: IGN Agent: Granite

Model	Release Date	Cut Off Date (Knowledge)	Token Costs (per Million tokens) and Context Window (CW)	Model Type	Usage Reasoning
Granite 4.0 Micro	Oct 20, 2025	Apr 2024	I/P: 0.017\$/M O/P: 0.11\$/M CW: 131K	Open Source	Model is fine-tuned for long context tool calling
Granite 4.1 8B	Apr 30, 2026	Early 2026	I/P: 0.05\$/M O/P: 0.10\$/M CW: 131K	Open Source	The model is designed for enterprise tasks including tool calling, retrieval-augmented generation (RAG), code generation with fill-in-the-middle

ASSIGNMENT 1
10 Different Model Providers and their Models released last 12 months

					support, text summarization, classification, and extraction.
--	--	--	--	--	---

8. Provider: Nvidia

Model	Release Date	Cut Off Date (Knowledge)	Token Costs (per Million tokens) and Context Window (CW)	Model Type	Usage Reasoning
Nemotron 3 Super	Mar 11, 2026	Dec 2025	I/P: 0.09\$/M O/P: 0.45\$/M CW: 263K - 1M	Open Source hybrid	The model features a 1M token context window for long-term agent coherence, cross-document reasoning, and multi-step task planning.
Nemotron 3 Nano 30B A3B	Dec 14, 2025	Nov 28, 2025	I/P: 0.05\$/M O/P: 0.20\$/M CW: 263K	Open Source hybrid	The model is small language MoE and fully open with open-weights, datasets and recipes so developers can easily customize, optimize, and deploy the model on their infrastructure for maximum privacy and security.
Nemotron 3 Nano Omni	Apr 28, 2026	Nov 28, 2025	I/P: 0\$/M (Free) O/P: 0\$/M(Free) CW: 256K	Open Source hybrid	Omni-modal reasoning model that processes text, audio, images, and video and produces text output, enabling agents to perceive and reason across modalities in a single inference loop.

ASSIGNMENT 1
10 Different Model Providers and their Models released last 12 months

9. Provider: xAI Model: Grok

Model	Release Date	Cut Off Date (Knowledge)	Token Costs (per Million tokens) and Context Window (CW)	Model Type	Usage Reasoning
Grok 3	Jun 10, 2025	Feb 28, 2025	I/P: 3\$/M O/P: 15\$/M CW: 131K	private Model	Grok 3 directly integrates with the X platform. This allows it to analyze breaking news, current trends, and real-time conversations, offering a distinct edge in timely information, dynamic deepsearch by scanning Internet & X to provide comprehensive, up-to-date answers for complex queries, offering a transparent view of how it arrived at its conclusions.
Grok 4	Jul 10, 2025	Jul 31, 2025	I/P: 3\$/M O/P: 15\$/M CW: 256K	private Model	Superior reasoning, a massive 256k context window, and native real-time integration with X
Grok 4.1 Fast	Nov 19, 2025	Nov, 2025	I/P: 0.20\$/M O/P: 0.50\$/M CW: 2M	private Model	Best model for real world use cases like customer support, deep search, etc.
Grok 4.2	Mar 31, 2026	Sep 1, 2025	I/P: 1.25\$/M O/P: 2.50\$/M CW: 2M	private Model	It provides the lowest hallucination rate on the market with strict prompt adherence, delivering consistently precise and truthful responses.
Grok 4.3	Apr 30, 2026	Dec, 2025	I/P: 1.25\$/M O/P: 2.50\$/M CW: 1M	private Model	A reasoning model that accepts text and image inputs with text output, and is suited for agentic workflows,

ASSIGNMENT 1
10 Different Model Providers and their Models released last 12 months

					instruction-following tasks, and applications requiring high factual accuracy.
--	--	--	--	--	---

10. Provider: Moonshot AI

Model	Release Date	Cut Off Date (Knowledge)	Token Costs (per Million tokens) and Context Window (CW)	Model Type	Usage Reasoning
Kimi K2.6	Apr 20, 2026	Jan 2025	I/P: 0.75\$/M O/P: 3.50\$/M CW: 262K	Open Source public	Optimized for long-horizon coding and agentic tasks, offering 12+ hour autonomous sessions and 300+ parallel agent swarms. It delivers state-of-the-art coding (76.8% SWE-Bench) and native multimodal capabilities (text/image/video) at significantly lower costs than proprietary models
Kimi K2.5	Jan 27, 2026	Jan 2025	I/P: 0.44\$/M O/P: 2\$/M CW: 262K	Open Source public	Native multimodal model, delivering state-of-the-art visual coding capability and a self-directed agent swarm paradigm.
Kimi K2 Thinking	Nov 6, 2025	Dec, 2024	I/P: 0.60\$/M O/P: 2.50\$/M CW: 262K	Open Source public	A model specializing in advanced, long-horizon agentic workflows, offering superior reasoning, deep tool integration, and high

ASSIGNMENT 1
10 Different Model Providers and their Models released last 12 months

					cost-efficiency compared to proprietary models. Key advantages include stable, multi-step planning with 200–300 consecutive tool calls and state-of-the-art performance on reasoning-heavy task
--	--	--	--	--	--